Machine Learning

1. Motivation + Theorie

Siegfried Gessulat

SAP Health Potsdam

Technical University of Munich
Chair of Proteomics and and Bioanalytics

s.gessulat@tum.de

FH Ludwigshafen
2018-10-08

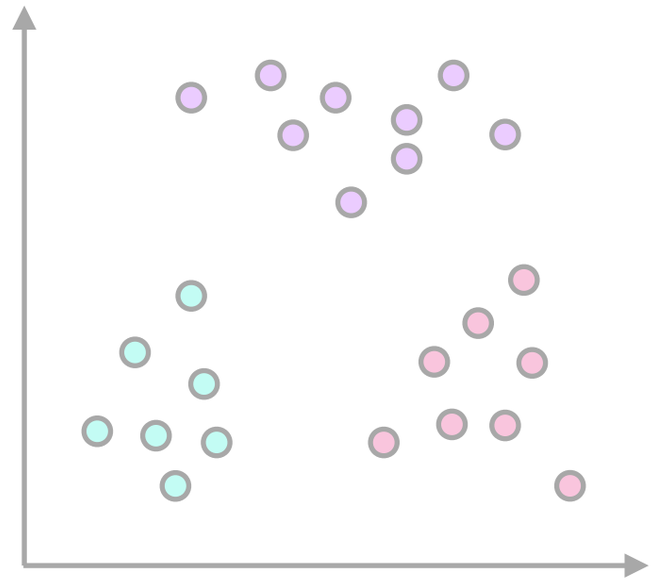

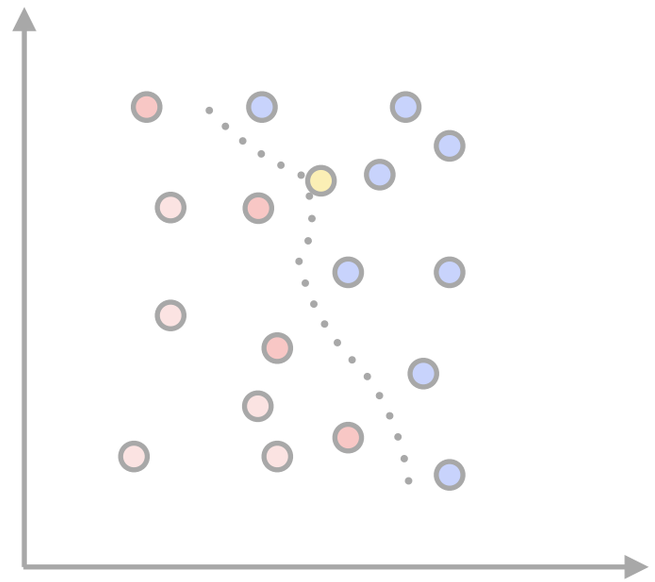

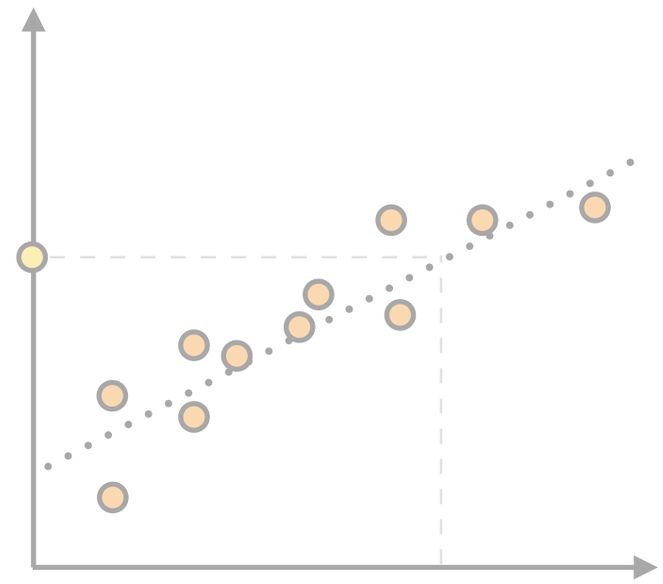

Course Outline

Block I Foundations

Oct 08: Introduction
- Overview machine learning
- Theory: Linear Algebra
- Algorithms: Knn, K-means

Oct 09: Basics
- Theory: linear regression, logistic regression
- Algorithms: gradient descent

Block II Best practices

Oct 29: Neural Networks
- Data cleaning
- Algorithm: Neural Networks

Oct 30: Best practices
- Theory: Cross validation
- Theory: Regularization

Course Outline

Block III Dark Arts

Nov 19: Tricks of the Trade
- Ensembles
- Hyperparameter Search
- Deep Learning Black Magic

Nov 20: Outlook
- Theory: Dimensionality Reduction

Outline Today

1. Preliminaries
2. Dataset: MNIST
3. What is Machine Learning?
4. Notation
5. Classification: K nearest neighbours (Knn)
6. Clustering: K-means
7. Theory: Linear Algebra refresher
8. Application: Python Intro, Implementation
9. Dataset: CIFAR-10

Preliminaries

1. Programming Assignments
Teams of two (1 SAP 1 non-SAP), randomly assigned.

2. Schedule, Dates & Deadlines
3 blocks. 1 assignment per block.
Help desk: Monday after each block - doodle your slot.
Deadlines: Friday after the help desk.
Results: Friday after the deadline.

Lecture Oct 8 **Mo**	Oct 9 **Tue**	**Wed**	**Thu**	Oct 12 **Fr**		
Help Oct 15 **Mo**	**Tue**	**Wed**	**Thu**	**Deadline** Oct 19 **Fr**		
Oct 22 **Mo**	**Tue**	**Wed**	**Thu**	**Results** Oct 26 **Fr**		

Preliminaries

1. Programming Assignments
Teams of two (1 SAP 1 non-SAP), randomly assigned.

2. Schedule, Dates & Deadlines
3 blocks. 1 assignment per block.
Help desk: Monday after each block
Deadlines: Friday after the help desk.
Results: Friday after the deadline.

What's your schedule Monday?

3. Grades
Assignment I: 30 points
Assignment II: 30 points
Assignment III: 40 points
Late assignments will cost you points.

4. Resources
After the end of one block, you will get an email with slides, code demonstrations, assignments and a doodle link to the Help desk.

Dataset: MNIST

http://yann.lecun.com/exdb/mnist/

70,000 samples
images of handwritten digits
28x28 grayscale images
labels are digits from 0-9

image: wikipedia.com

image: research.fb.com

Yann LeCun

NYU Professor
facebook Chief AI Scientist
Deep Learning (ConvNets)

image: di.ku.dk

Corinna Cortes

Head Google Research NY
Support Vector Machines

image: microsoft.com

Christopher J.C. Burges

formerly Microsoft Research
Support Vector Machines

What is Machine Learning?

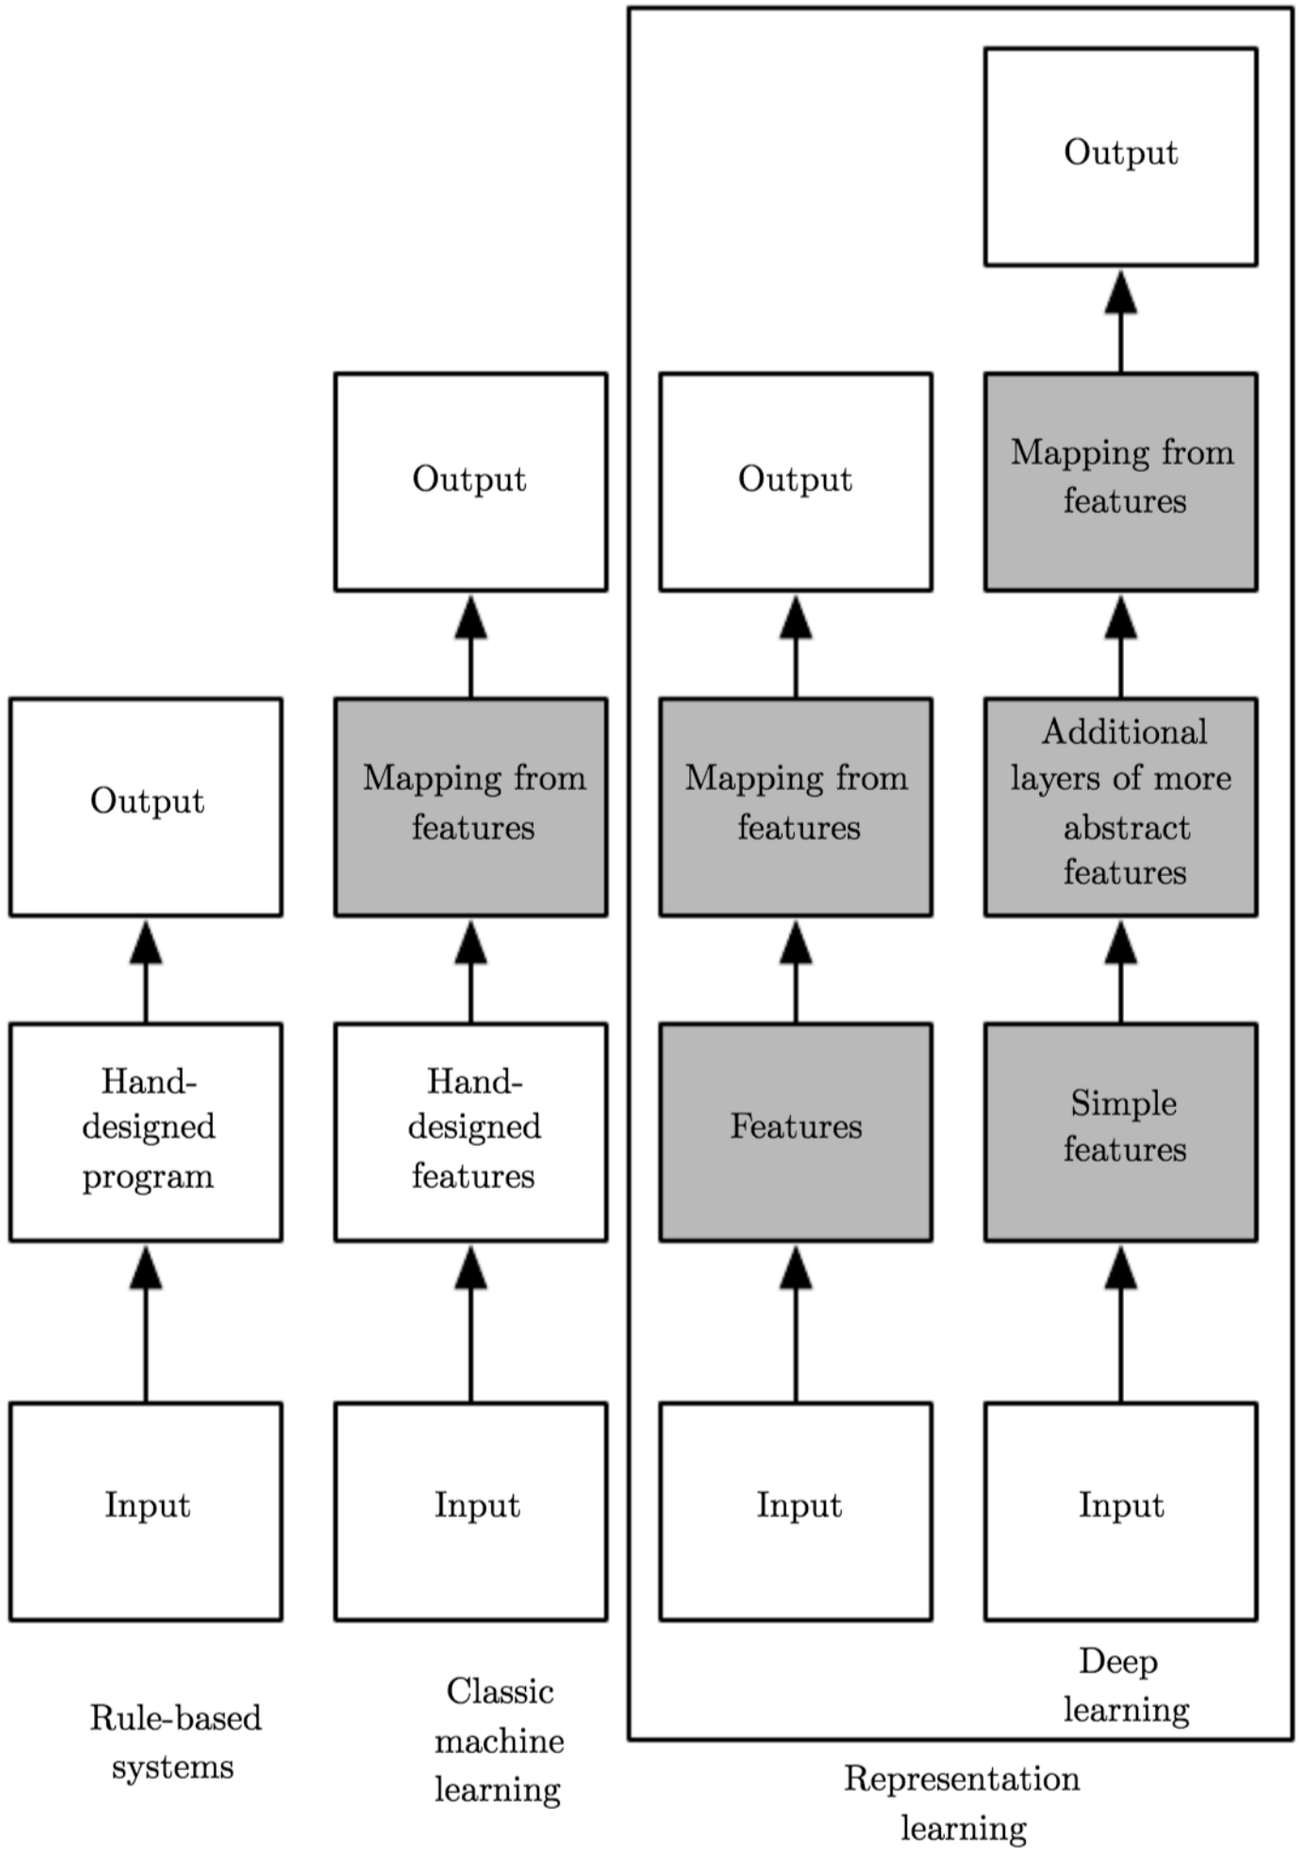

graphic: "Deep Learning" Goodfellow, Bengio, Courville

Machine Learning Overview

Do you have labeled data?

✓ supervised

what kind of label?

continuous → regression

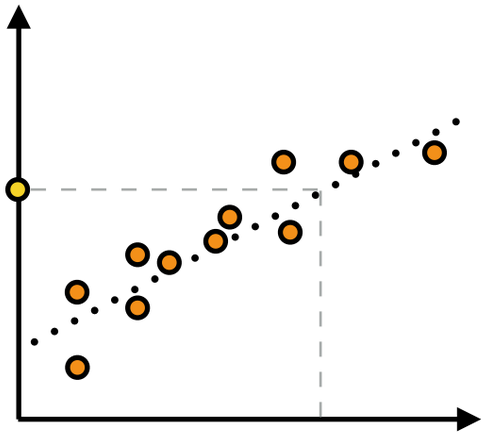

Boston Housing

predict real estate price by attributes of the property

discrete → classification

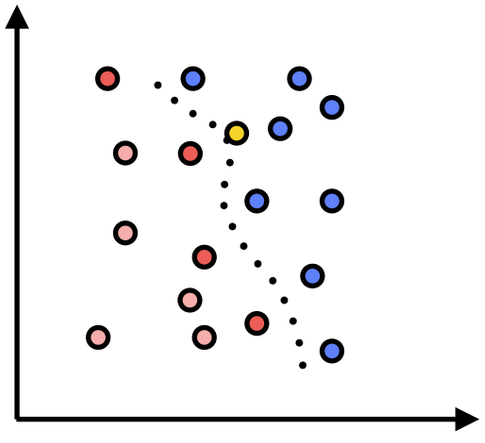

MNIST

classify handwritten digit as 0-9

✗ unsupervised

what to generate?

sample → label → clustering

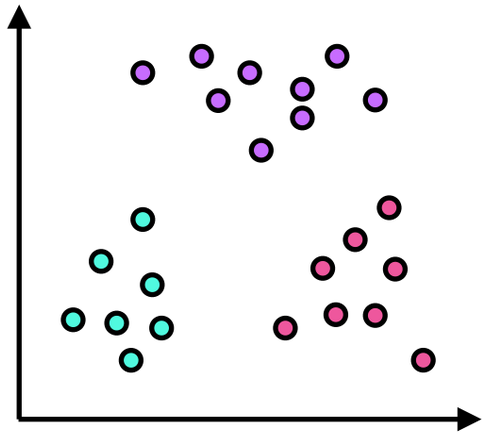

MNIST

cluster handwritten digits in 10 clusters

sample → sample → generative

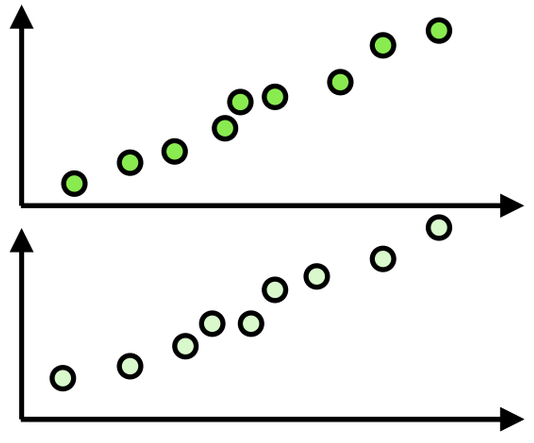

MNIST

generate images of handwritten digits

Machine Learning Overview

Do you have labeled data?

✓ → **supervised**
what kind of label?

- continuous → **regression**

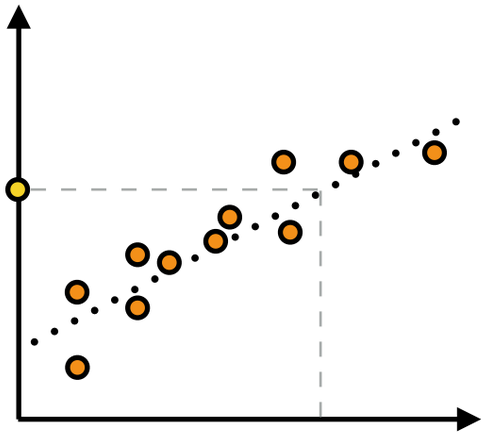

- K-nearest-neighbours
- Linear Regression
- Regression Trees
- Support Vector Regression
- Neural Networks

- discrete → **classification**

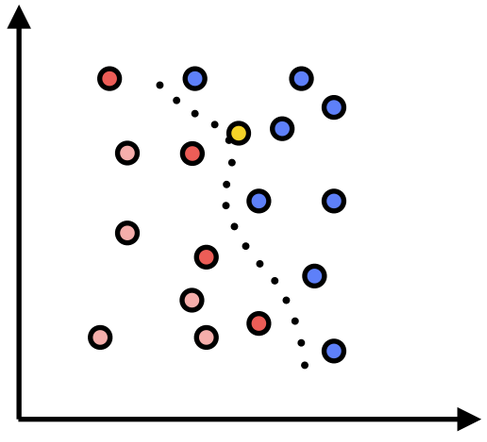

- Logistic Regression
- Support Vector Machines
- Decision Trees / Forests
- Neural Networks
- (Gaussian) Mixture Model

✗ → **unsupervised**
what to generate?

- sample → <u>label</u> → **clustering**

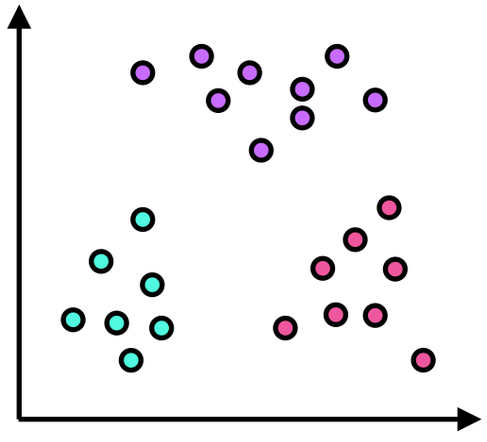

- K-means
- (Gaussian) Mixture Models
- Self-organising maps*

- sample → <u>sample</u> → **generative**

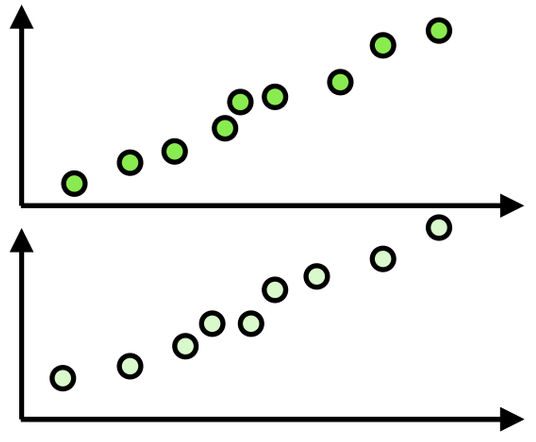

- Markov Chains
- Autoencoder*
- Generative Adversarial Networks*

* neural network

Machine Learning Overview

Do you have labeled data?

✓ → **supervised**
what kind of label?

- continuous → **regression**

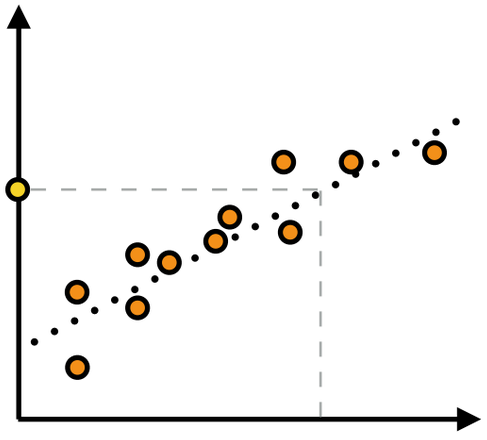

- K-nearest-neighbours
- Linear Regression
- Regression Trees
- Support Vector Regression
- Neural Networks

- discrete → **classification**

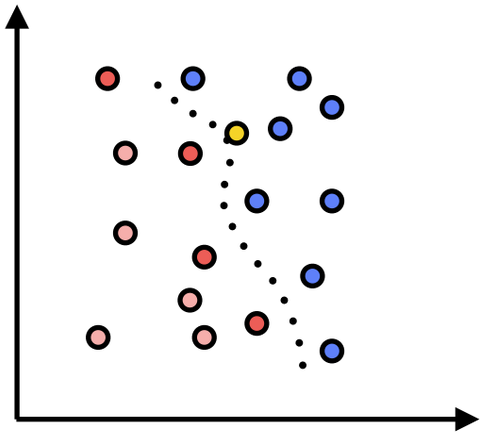

- Logistic Regression
- Support Vector Machines
- Decision Trees / Forests
- Neural Networks
- (Gaussian) Mixture Model

✗ → **unsupervised**
what to generate?

- sample → <u>label</u> → **clustering**

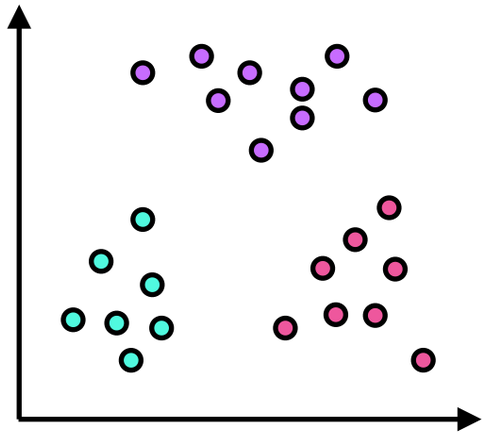

- K-means
- (Gaussian) Mixture Models
- Self-organising maps*

- sample → <u>sample</u> → **generative**

not covered

- Markov Chains
- Autoencoder*
- Generative Adversarial Networks*

* neural network

Machine Learning Overview

Do you have labeled data?

✓ → **supervised**

what kind of label?

- continuous → **regression**

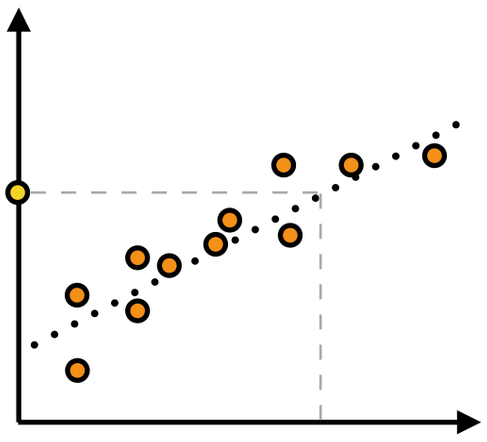

- K-nearest-neighbours

Oct 08

- Linear Regression

Oct 09

- discrete → **classification**

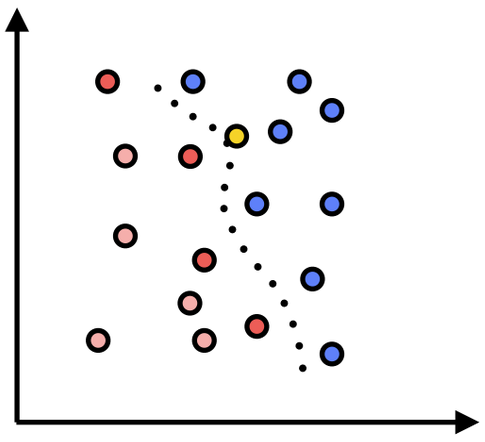

- Logistic Regression
- Support Vector Machines

Oct 09

- (Gaussian) Mixture Model

✗ → **unsupervised**

what to generate?

- sample → label → **clustering**

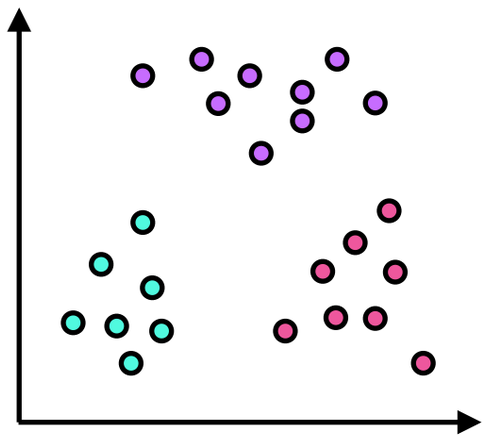

- K-means

Oct 08

- Self-organising maps*

- sample → sample → **generative**

not covered

- Markov Chains
- Autoencoder*
- Generative Adversarial Networks*

* neural network

Notation

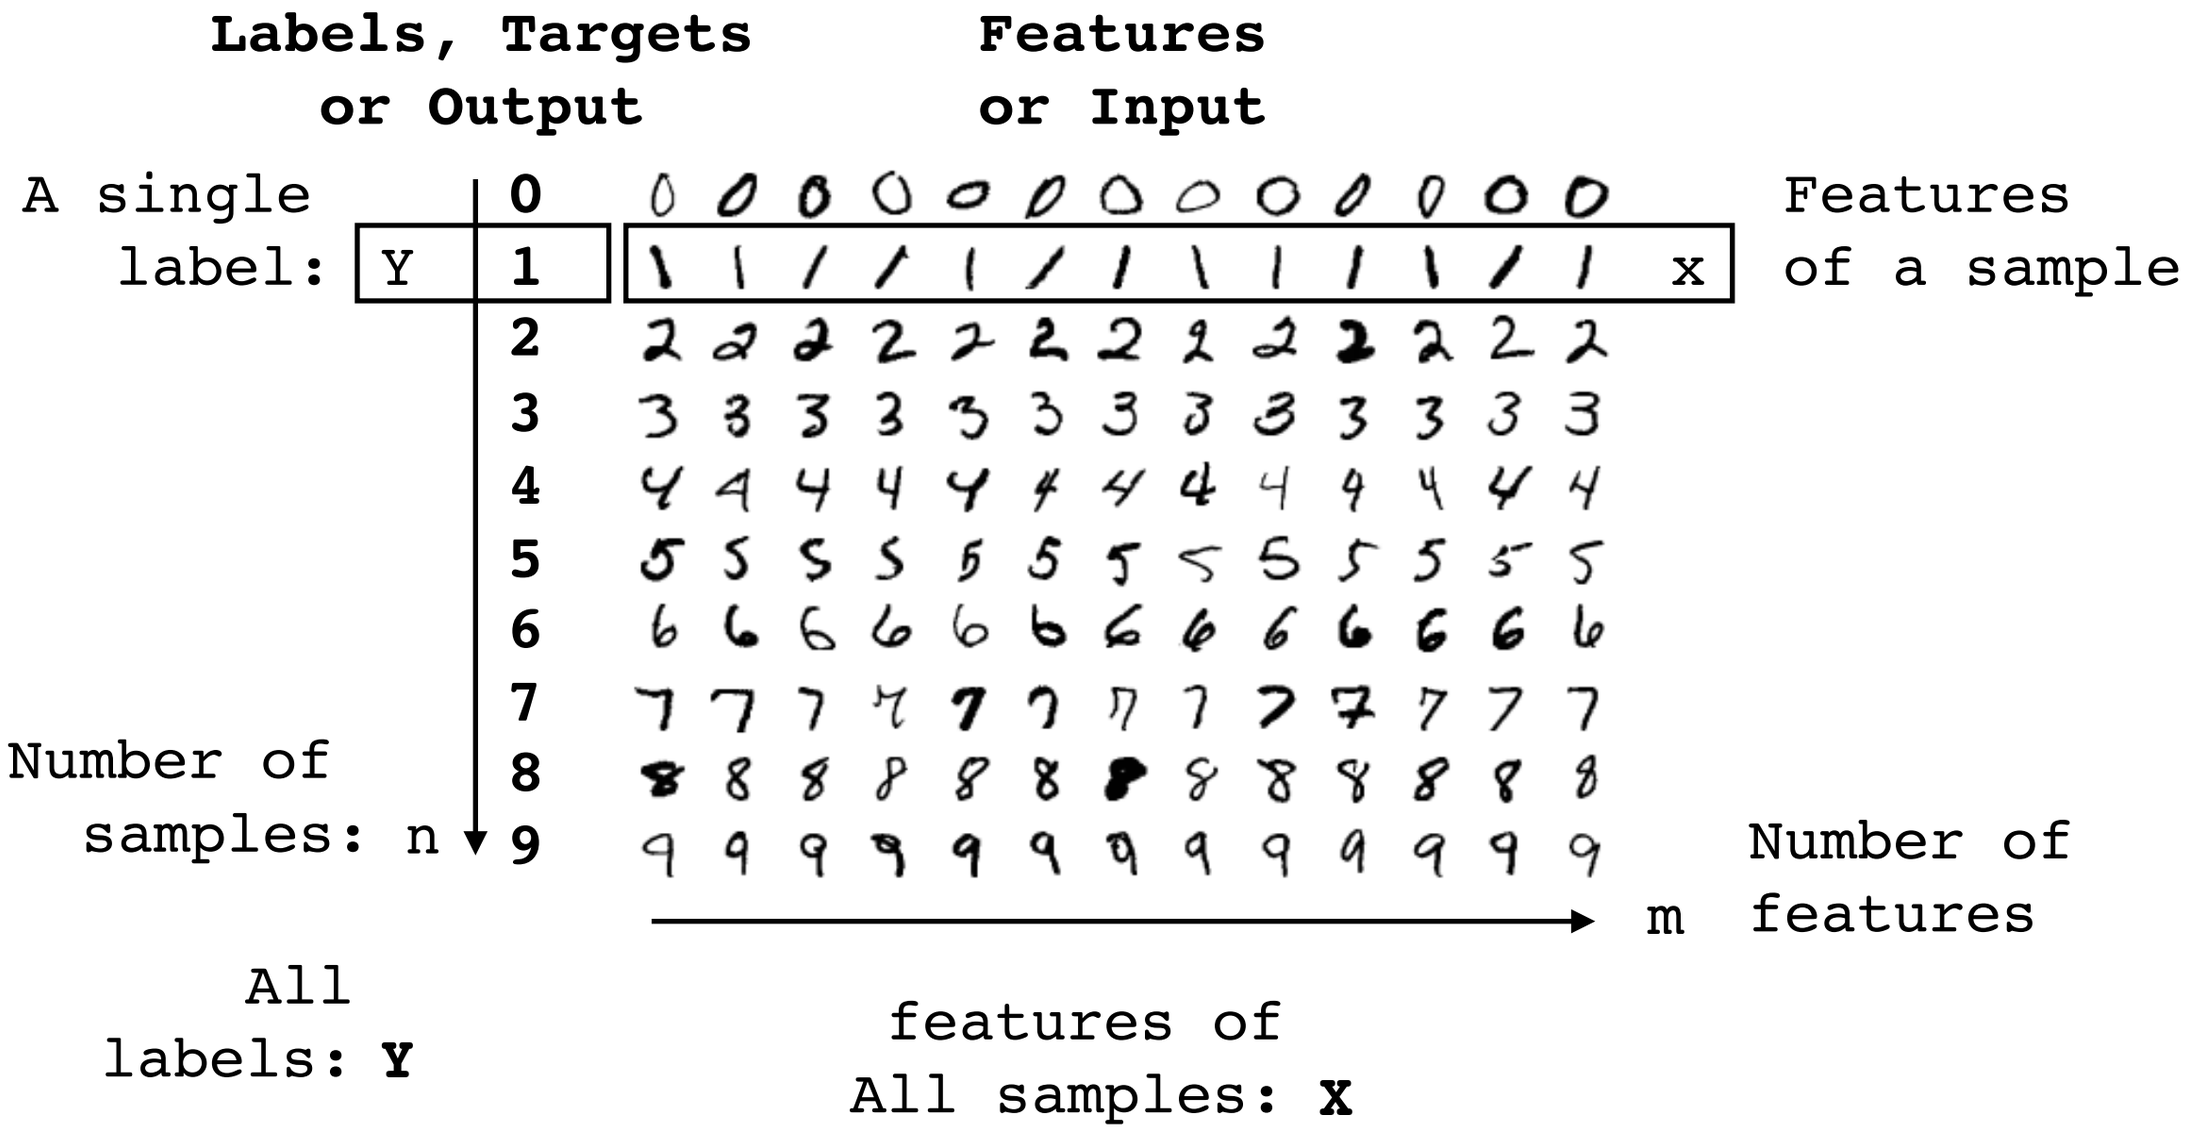

A machine learning model is a function mapping X to Y.
Θ is what the model “learned”.

$$f_\theta : X \rightarrow Y$$

Classification: k-nearest neighbors (knn)

Intuition

1. Of all neighbors, which are the k nearest to my sample?

2. What label does the majority of the k neighbors

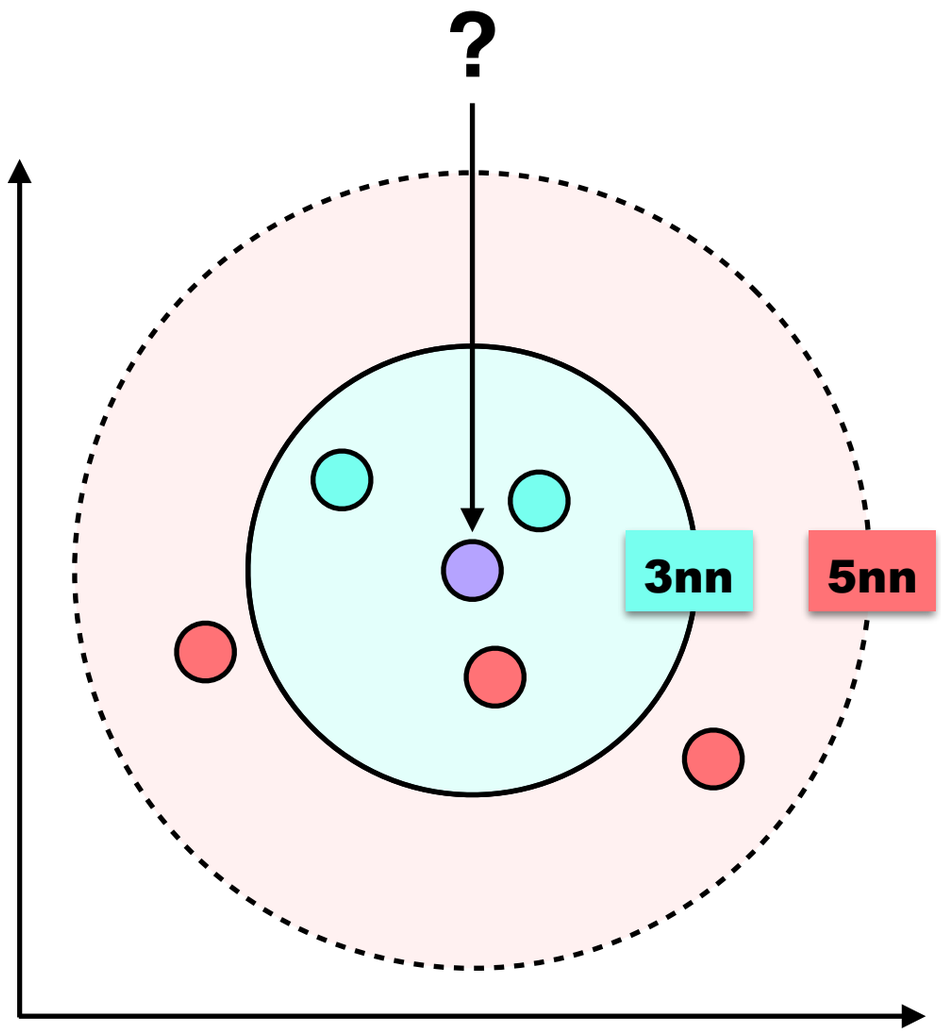

Classification: k-nearest neighbors (knn)

Algorithm

Input
- dataset **M** - matrix of samples (vector, class)
- sample **v** to classify
- number of neighbors **n**

Output
- class **k**

0. Calculate distance of **v** to all samples in **M**
1. Select the **n** samples closest to **v**
2. choose **k** to be the majority of classes in selection

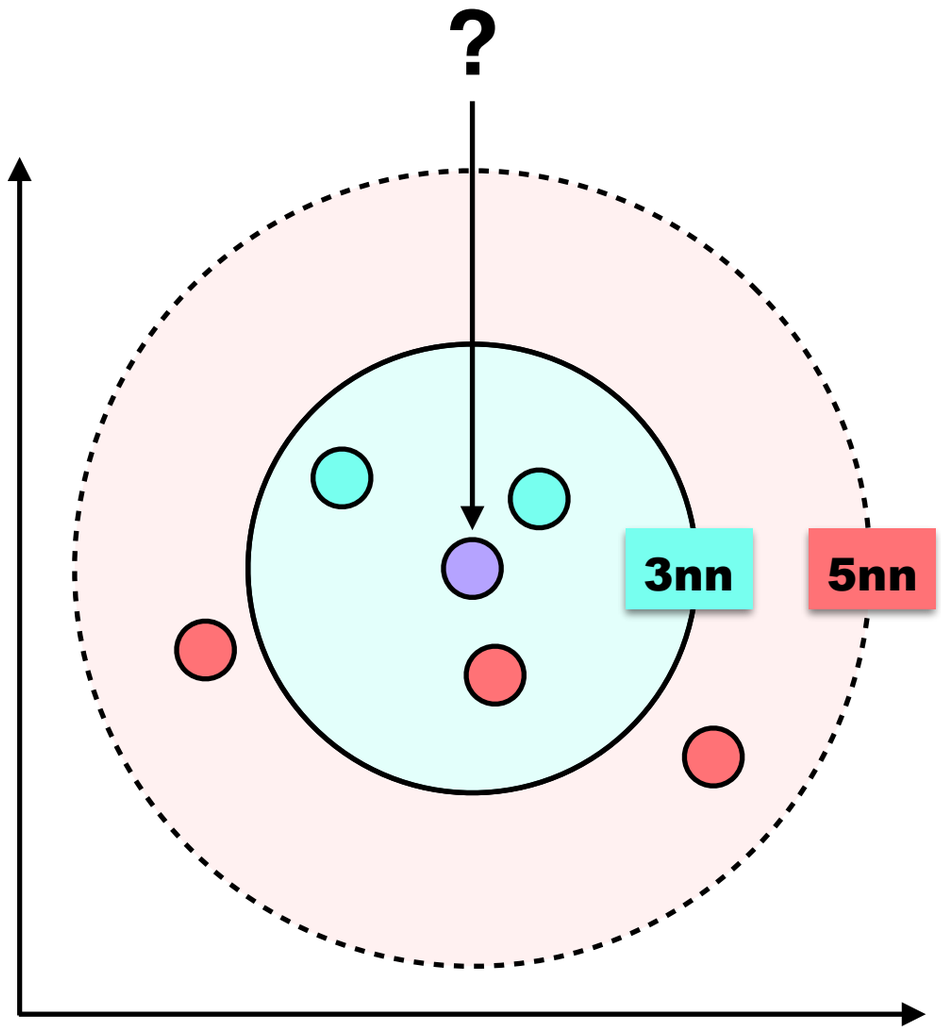

Classification: k-nearest neighbors (knn)

Data

2

?

5

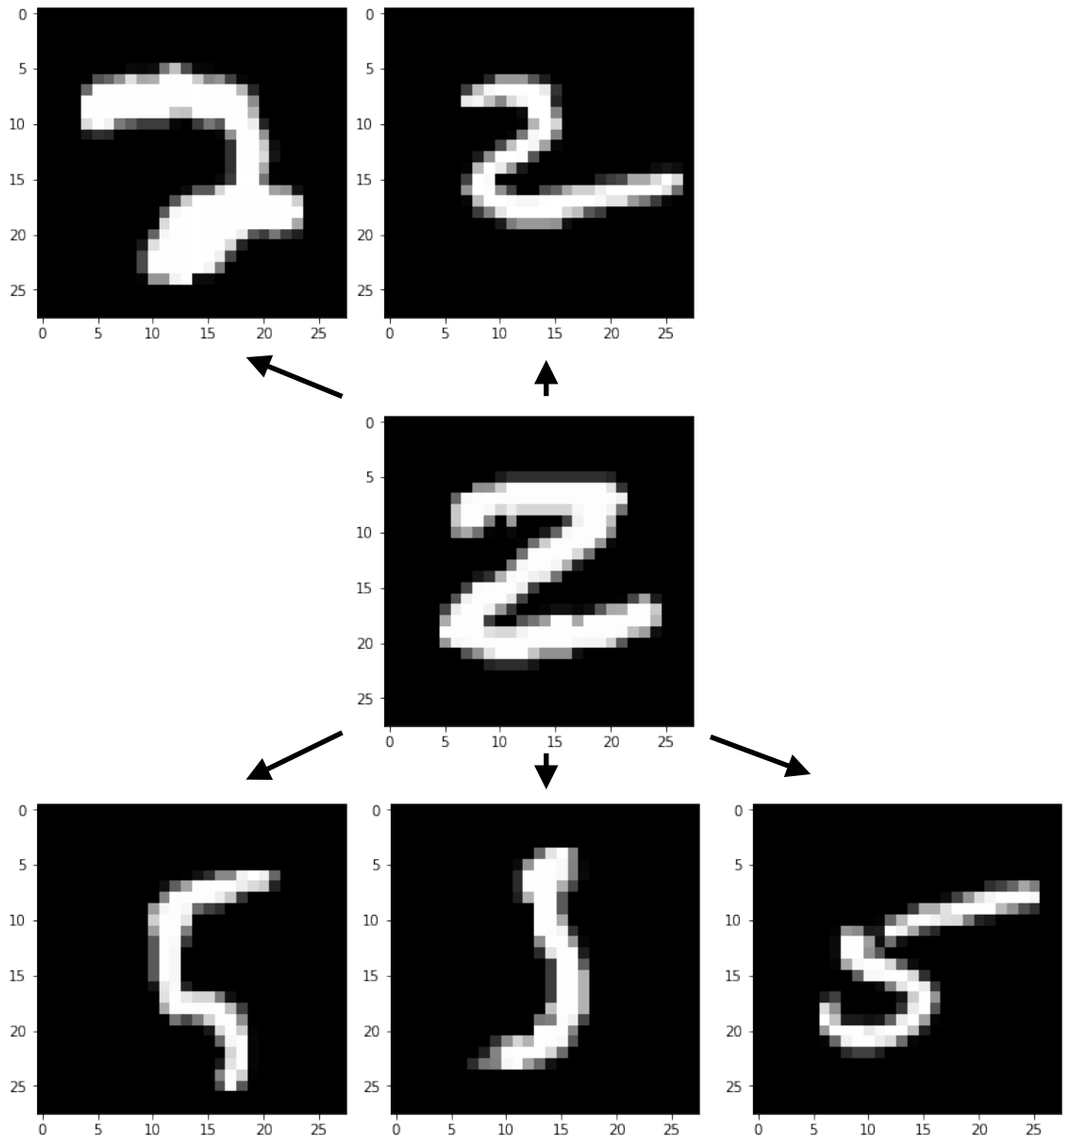

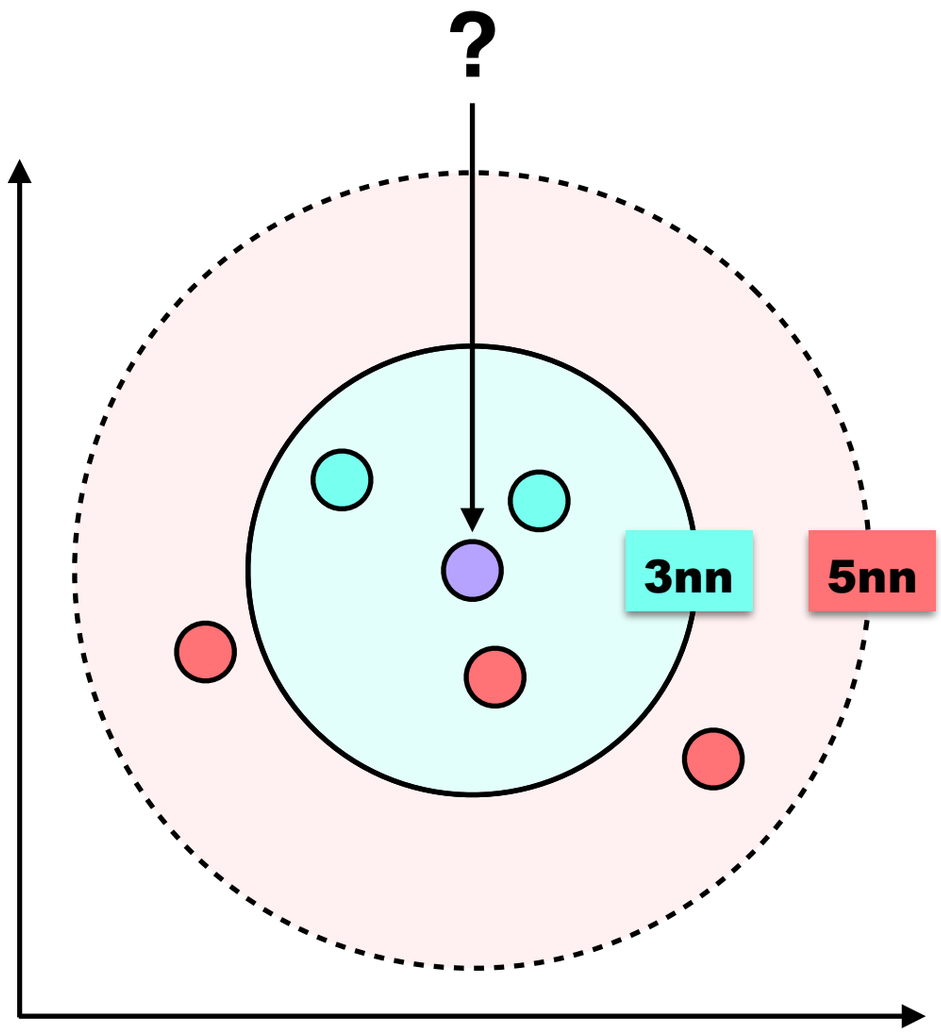

Clustering: k-means

Intuition

Initially, select k random samples, those are your class centroids.

1. assign each sample the class of the closest centroid.
2. determine centroid for each class with smallest distance to each all samples in class
3. if one of the centroids changed, repeat.

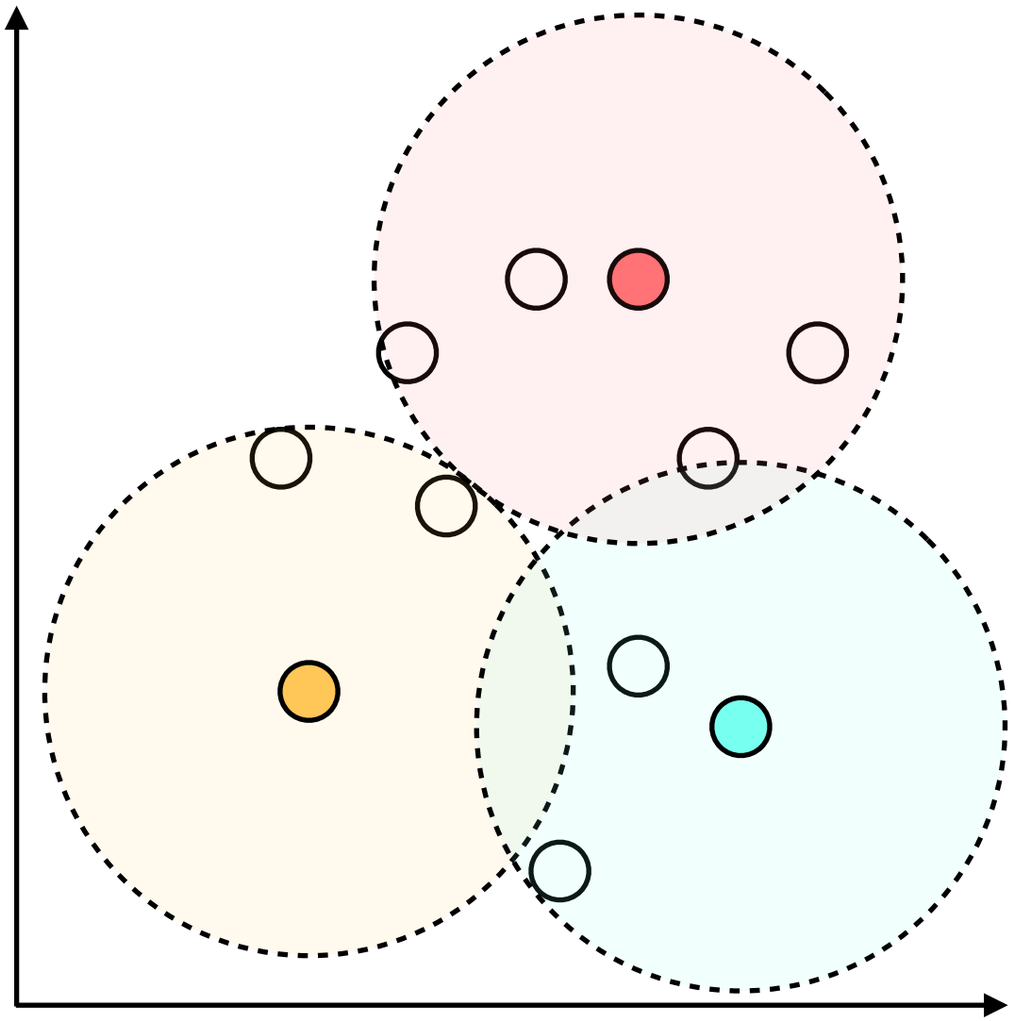

Clustering: k-means

Algorithm

Input
- dataset **M** - matrix of samples (vector)
- number of clusters **k**

Output
- dataset **R** - matrix of samples (vector, class)

1. select **k** samples from **M** at random (**K** cluster centroids)
2. assign all samples in **M** to the nearest sample in **K**
3. Calculate the cluster centroid **k'** for each cluster **k** in **K** as **K'.**
4. if **K != K'** then set **K** := **K'**; Repeat from 3.

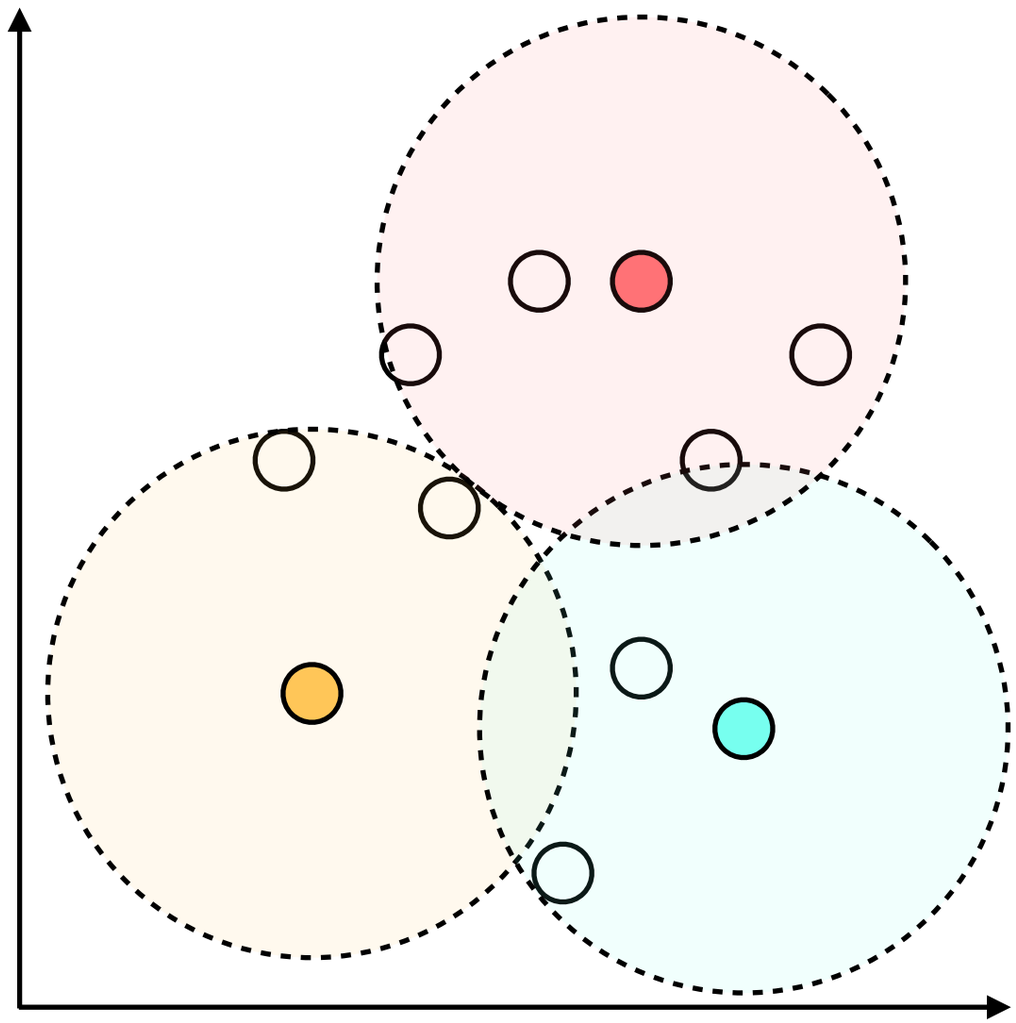

Clustering: k-means

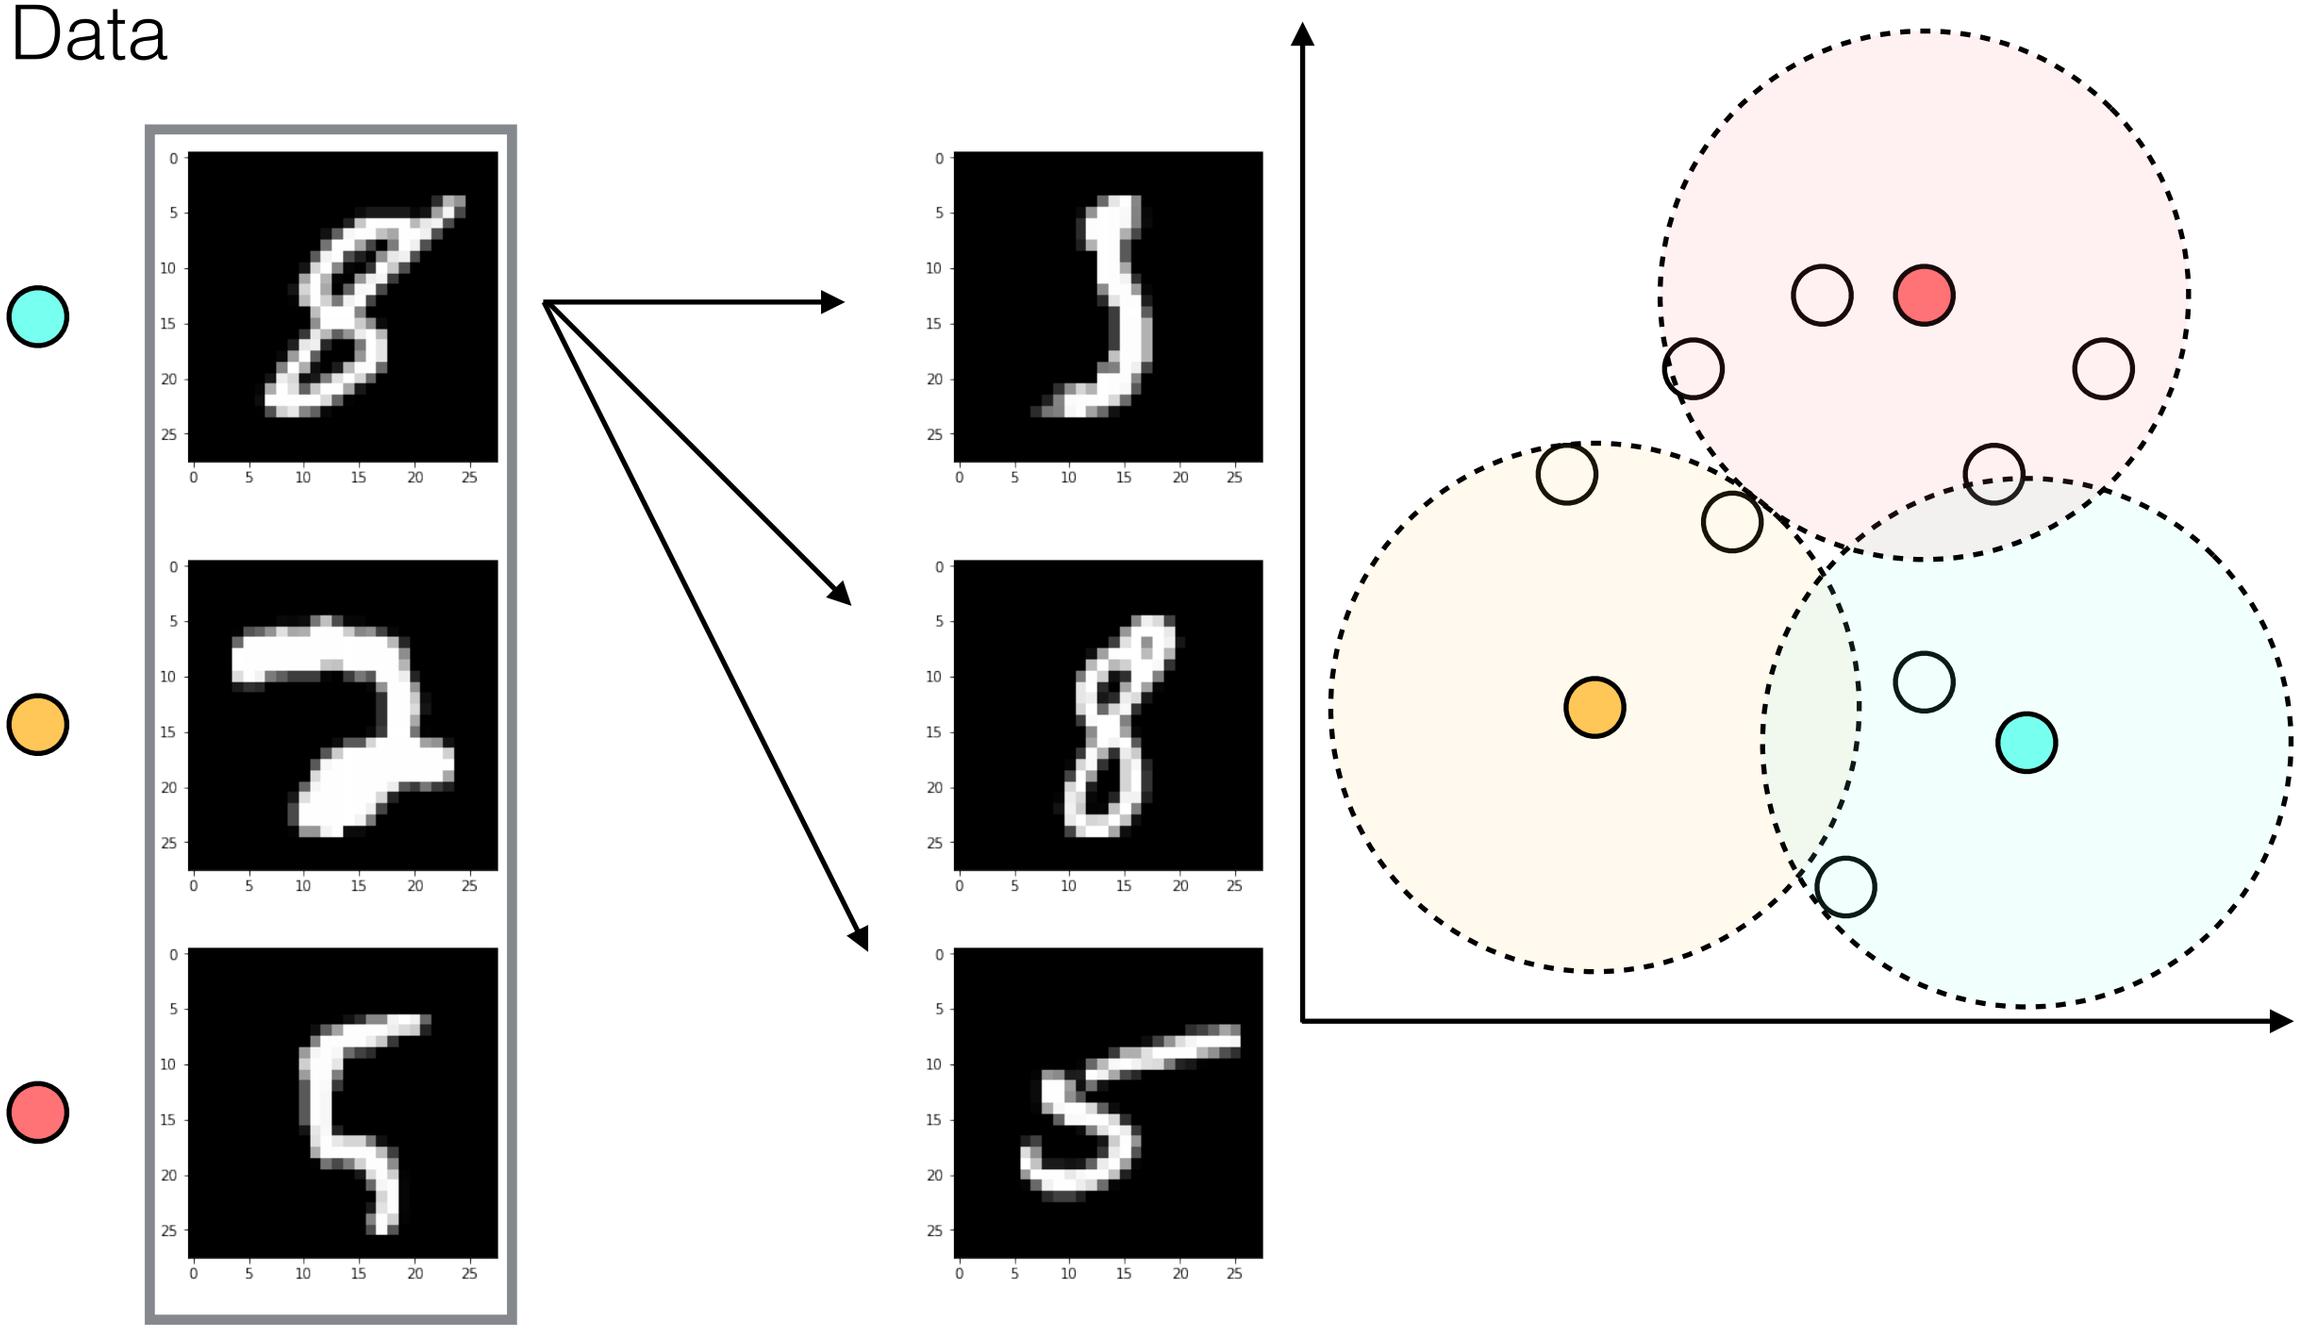

Theory: Linear Algebra

Euclidean distance

$$dist(x, y) = \|x - y\|_2 = \sqrt{\sum_{i=1}^{n}(x_i - y_i)^2}$$

Example

$$\sqrt{\sum_{i=1}^{2}\left(\begin{bmatrix}2\\3\end{bmatrix} - \begin{bmatrix}4\\1\end{bmatrix}\right)^2} = \sqrt{\sum_{i=1}^{2}\left(\begin{bmatrix}-2\\2\end{bmatrix}\right)^2}$$

$$= \sqrt{\sum_{i=1}^{2}\begin{bmatrix}4\\4\end{bmatrix}}$$

$$= \sqrt{8}$$

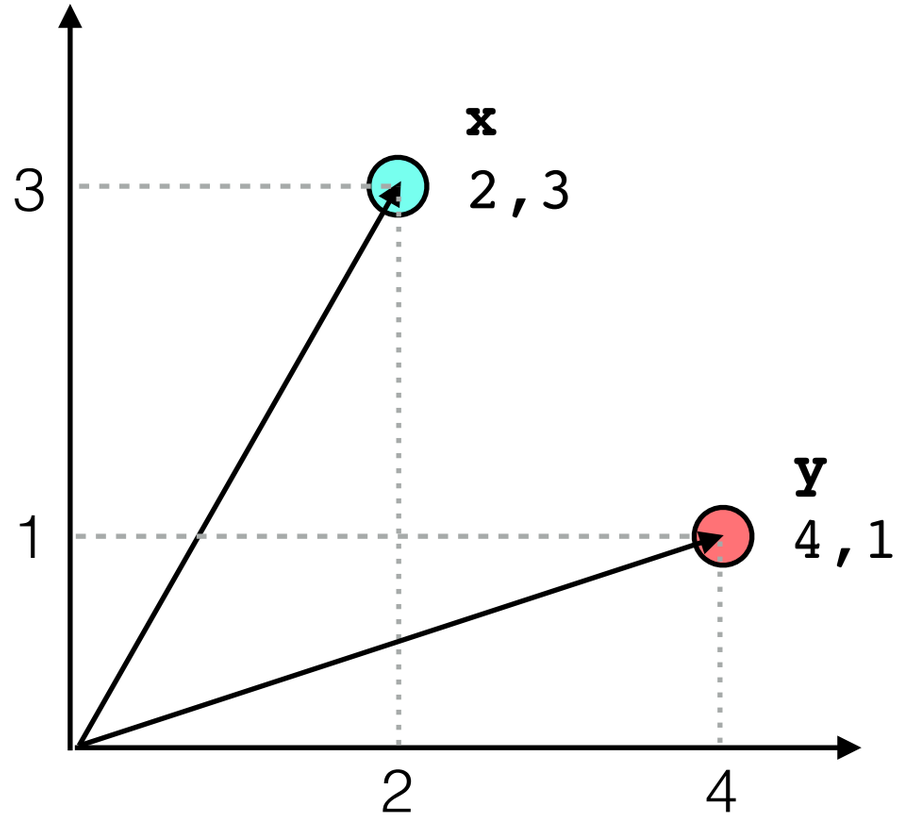

Code

```
dist = np.linalg.norm(x - y)
```

Dataset: CIFAR-10

cs.toronto.edu/~kriz/cifar.html

60,000 samples
images of objects
32x32 color images with labels
labels are 10 categories

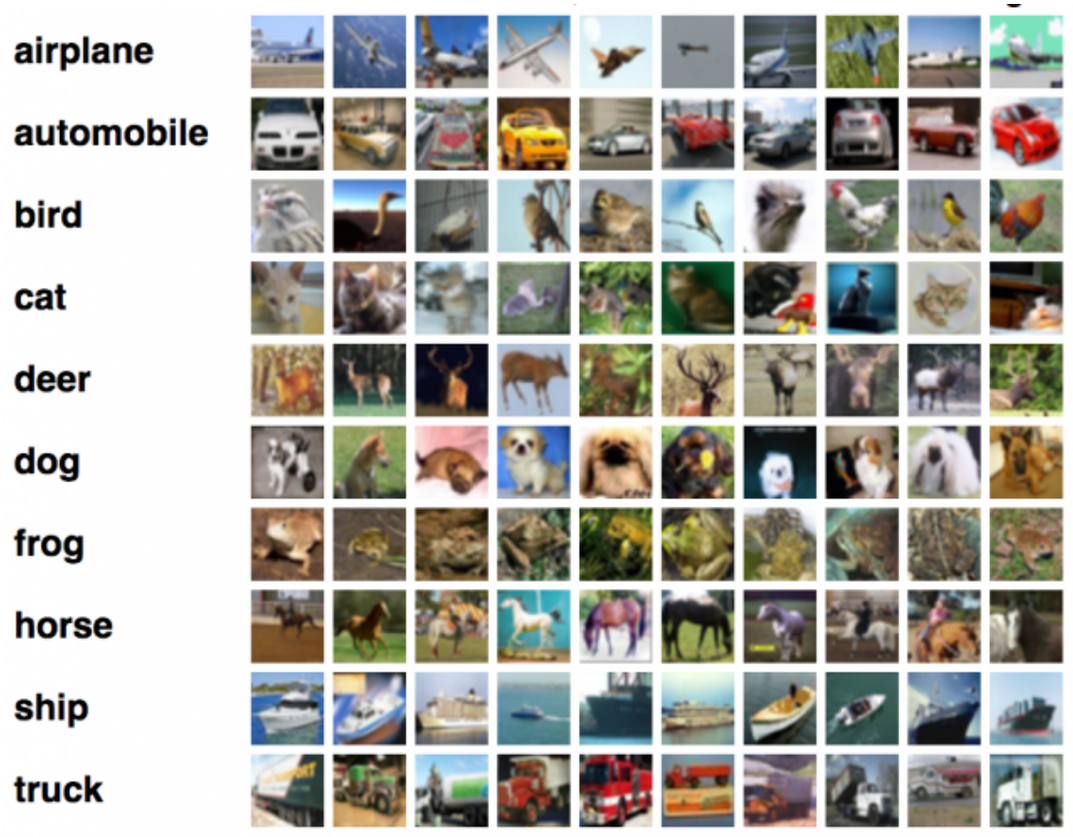

image: kaggle.com

image: qz.com

Alex Krizhevsky

Dessa, formerly Google
Deep Learning (AlexNet)

image: cs.toronto.edu

Vinod Nair

Yahoo Labs
Deep Learning

image: thestar.com

Geoffrey Hinton

U Toronto Professor
Google Brain
Deep Learning